

AU 332 Artificial Intelligence: Principles and Techniques

Homework 3

Due OCT 31st 11:59pm

Adhere to the Code of Academic Integrity. You may discuss background issues and general strategies with others and seek help from course staff, but the implementations that you submit must be your own. In particular, you may discuss general ideas with others but you may not work out the detailed solutions with others. It is never OK for you to see or hear another student's code and it is never OK to copy code from published/Internet sources. If you feel that you cannot complete the assignment on your own, seek help from the course staff.

When submitting your assignment, follow the instructions summarized in Section 4 of this document.

1 Reinforcement Learning in Maze Environment

In this assignment, you will implement a Dyna-Q learning agent to search for the treasure and exit in a grid-shaped maze. The agent will learn by trial and error from interactions with the environment and finally acquire a policy to get as high as possible scores in the game. Please finish this part *individually*.

1.1 Game Description

Suppose a 6×6 grid-shaped maze in Figure 1. The red rectangle represents the start point and the green circle represents the exit point. You can move upward, downward, leftward and rightward and you should avoid falling into the traps, which are represented by the black rectangles. Finding the exit will give a reward +1 and falling into traps will cause a reward -1, and both of the two cases will terminate current iteration. Moreover, you will get a bonus reward +2 if you find the treasure, which shown as golden diamond.




Figure 1: Maze environment

The environments are implemented in `maze_env1.py` and `maze_env2.py`. The state space and action space are briefly described as follows and you may learn more details in the code.

State(5-dimension): The current position of the agent (4D) and a bool variable (1D) that indicates whether the treasure has been found.

Action(1-dimension): A discrete variable and 0, 1, 2, 3 respectively represent move upward, downward, rightward and leftward.

Note: Please *do not* revise the `maze_env1.py` and `maze_env2.py`.

1.2 Dyna-Q Learning [50 points]

You are asked to implement an agent based on Dyna-Q learning, which can find optimal policy both in `maze1` and `maze2` in Figure 1. Please design the agent in `agent.py` and complete the training process in `main.py`. An agent with random policy is provided in `main.py`. You can learn how to interact with the environment through the demo and then write your own code.

1.3 Exploration-Exploitation Strategy [20 points]

Considering the explore-exploit dilemma, complete your Dyna-Q learning agent by **implementing an exploration-exploitation strategy**, meaning it chooses random actions in a fraction of the time to explore better policy, and follows its current best Q-values otherwise. Note that choosing a random action should not choose a random sub-optimal action, but **rather any random legal action**.

You should balance the exploration with exploitation. In the early stage of the training, exploration should be encouraged as we hope that the agent can find better policy (e.g. getting the treasure). Then with the advancement of training, exploration may be reduced for the divergence of the policy. **Our goal is to get high score with as few as possible interactions.**

There are many exploration strategies in reinforcement learning, you can read this paper and choose a suitable method for the agent: <https://www.cs.mcgill.ca/~cs526/roger.pdf>

2 Reinforcement Learning on Atari Game

In this part, you will implement a DQN agent to play atari game *breakout*. The given code is almost complete. You are asked to read the code and complete the remaining part. The performance of the agent may not be satisfactory and you have to tune it to get higher scores. You will finish this part **in a group of two**.

2.1 Game Description

Atari breakout is one of the classic and leading games. You need to bounce back the ball to crush all of the boxes which are showing upward. In this assignment, you are asked to design a DQN agent to learn control policies directly from the visual information of the game.




Figure 2: Atari Breakout

2.2 Deep Q-Learning [20 points]

Considering the complexity of the task and the limitation of hardware requirements, you will implement a simple DQN agent in this part and set a limitation on the number of iterations. You need to understand the code and complete the training process in **main()**. Before running your code, you need to install the required environment for it (see Section 3).

2.3 Tune the Agent [10 points]

The agent in Section 2.2 may not get a high score on the game. You need to tune the agent to get as high as possible scores within 5000 iterations of the training. Here are some possible ways to improve the performance:

- Tuning the hyper-parameters of the agent.
- Tuning the structure of the network.
- Implementing experience replay.
- Your own ideas.

Please describe your method and the tuning process in your report.

3 Installation

To train the agent on atari game, you need to install Gym and Keras. You can follow the tutorial in this section to install the environment on Linux, Windows or macOS, and we strongly recommend you use Linux system.

3.1 Install Anaconda

Open the address <https://www.anaconda.com/distribution/> and download the installer of **Python 3.7 version** for your system.

- For Linux
bash Downloads/Anaconda3-2019.07-Linux-x86_64.sh
- For Windows
Open the .exe file and follow the installer steps.
- For macOS
Open the .pkg file and follow the installer steps.

3.2 Install Required Environment

After installing anaconda, open a terminal (Linux and macOS) or Anaconda Prompt (Windows) and create an environment for Gym:

conda create python=3.6 --name gym

Then activate the environment

- For Linux and macOS: **source activate gym**
- For Windows: **activate gym**

Install gym

pip install gym

pip install gym[atari]

Install tensorflow and keras

pip install tensorflow

pip install keras

Install pandas

pip install pandas

Execute the code

cd

python atariDQN.py

4 Submission instructions

1. Zip all your python files **main.py**, **agent.py**, **atariDQN.py** and report file **HW3.pdf** to a folder called **homework3_name.zip**.
2. Send the zip file to TA *mingsun@sjtu.edu.cn*.